


Distributed Memory Programming using MPI

Kameswararao Anupindi

Department of Mechanical Engineering,
IIT Madras
February 2024

Example of a distributed memory machine



Credit: <https://www.nvidia.com/en-in/networking/infiniband-switching/>

Interconnects ~ 100 Gb/s

1. InfiniBand
2. Omni-Path

About MPI (Message Passing Interface)

- MPI 4.1 standard <https://www.mpi-forum.org/docs/>
- MPI Implementations:
 - OpenMPI, MPICH, MVAPICH, Intel MPI etc.
- MPI – not a new language, rather a library that provides functions for C/C++/FORTRAN.
- What is a process?
- Send - Recv function calls
- Point-to-point or collective communication
- Incremental parallelization not possible!
- SIMD/SPMD model

Compilation and Execution

- **mpicc** -g -Wall -o mpi_hello mpi_hello.c
- **mpif90** -g -Wall -o mpi_hello mpi_hello.f90
- mpicc/ mpif90 is a wrapper script (?)
- **mpiexec** -n <no. of processes> <executable>
- **mpiexec** -n 4 ./mpi_hello
- How do we get from mpiexec to execution of code?

Hello World - A serial program

```
/* Hello world program */  
#include <stdio.h>  
  
int main(void)  
{  
    printf("\n Hello, world!\n");  
  
    return 0;  
}
```

Hello World – Parallel version in C

```
#include<stdio.h>
#include<string.h>
#include<mpi.h>

int main(int argc, char** argv)
{
    int i, myid, size, tag=100;
    char message_send[50], message_recv[50];
    MPI_Status status;

    MPI_Init(&argc, &argv);

    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);

    /* printf("%d", size); */
}
```

Hello World – Parallel version in C contd.

```
if (myid != 0)
{
    sprintf(message_send, " Hello from process id: %d \n ", myid);
    MPI_Send(message_send, 50, MPI_CHAR, 0, tag, MPI_COMM_WORLD);
}
else
{
    for (i = 1; i < size; i++)
    {
        MPI_Recv(message_recv, 50, MPI_CHAR, i, tag, MPI_COMM_WORLD, &status);
        printf("\n %s", message_recv);
    }

    sprintf(message_send, " Hello from process id: %d \n ", myid);
    printf("\n %s", message_send);
}

MPI_Finalize();

return 0;
}
```

Hello World – Parallel version in FORTRAN

```
program main
  implicit none
  include "mpif.h"
  integer :: i, myid, size, mpierror, tag, status(MPI_STATUS_SIZE)
  character(len=50) :: message_send, message_recv

  call MPI_INIT(mpierror)

  call MPI_COMM_SIZE(MPI_COMM_WORLD, size, mpierror)
  call MPI_COMM_RANK(MPI_COMM_WORLD, myid, mpierror)

  tag = 100

  if (myid /= 0) then
    write(message_send, *), 'Hello from process id:', myid
    call MPI_SEND(message_send, 50, MPI_CHARACTER, 0, tag, MPI_COMM_WORLD, mpierror)
  else
    do i = 1, size-1
      call MPI_RECV(message_recv, 50, MPI_CHARACTER, i, tag, MPI_COMM_WORLD, status, mpierror)
      write(*, *), message_recv
    end do
    write(message_send, *), 'Hello from process id:', myid
    write(*, *), message_send
  end if

  call MPI_FINALIZE(mpierror)

end program main
```


Initialization and Finalization

```
int MPI_Init(int * argc_p, char *** argv_p)
```

```
int MPI_Init( NULL, NULL)
```

```
int MPI_Finalize(void)
```

```
subroutine MPI_INIT(mpierror)  
  integer, intent(out) :: mpierror
```

```
subroutine MPI_FINALIZE(mpierror)  
  integer, intent(out) :: mpierror
```

Communication

```
int MPI_Comm_size(  
    MPI_Comm comm,  
    int* comm_size_p);
```

```
int MPI_Comm_rank(  
    MPI_Comm comm,  
    int* myid_p);
```

```
subroutine MPI_COMM_SIZE(  
    MPI_Comm comm,  
    integer size,  
    integer mpierror)
```

```
subroutine MPI_COMM_RANK(  
    MPI_Comm comm,  
    integer myid,  
    integer mpierror)
```

Communication – C functions

```
int MPI_Send(  
    void*          message_buf_p,  
    int            message_size,  
    MPI_Datatype    message_type,  
    int            dest,  
    int            tag,  
    MPI_Comm        communicator);
```

```
int MPI_Recv(  
    void*          message_buf_p,  
    int            buf_size,  
    MPI_Datatype    buf_type,  
    int            source,  
    int            tag,  
    MPI_Comm        communicator,  
    MPI_Status*     status_p);
```

Communication – FORTRAN routines

```
int MPI_SEND(buf, count, datatype, dest, tag, comm, ierror)
    type(*), dimension(:), intent(in) :: buf
    integer, intent(in) :: count
    type(mpi_datatype), intent(in) :: datatype
    integer, intent(in) :: dest, tag
    type(mpi_comm), intent(in) :: comm
    integer, optional, intent(out) :: ierror
```

```
int MPI_RECV(buf, count, datatype, source, tag, comm, status,
             ierror)
    type(*), dimension(:) :: buf
    integer, intent(in) :: count
    type(mpi_datatype), intent(in) :: datatype
    integer, intent(in) :: source, tag
    type(mpi_comm), intent(in) :: comm
    type(mpi_status) :: status
    integer, optional, intent(out) :: ierror
```

Message Matching

Process a:

```
MPI_Send(send_buf_p, send_buf_size, send_type,  
         dest, send_tag, send_comm);
```

Process b:

```
MPI_Recv(recv_buf_p, recv_buf_size, recv_type,  
         src, recv_tag, recv_comm, &status)
```

recv_comm	==	send_comm
recv_tag	==	send_tag
dest	==	b
src	==	a

Message Matching contd.

Must specify compatible buffers:

send_buf_p	--	recv_buf_p
send_buf_size	--	recv_buf_size
send_type	--	recv_type

Most of the time, it is sufficient to have:

recv_type	==	send_type
recv_buf_size	>=	send_buf_size

Then the message sent by Process a can be successfully received by Process b.

The case of unknown message order

```
for ( i = 1; i < comm_size ; i++)  
{  
    MPI_Recv(recv_buf_p, recv_buf_size, recv_type,  
    MPI_ANY_SOURCE, recv_tag, comm, status_p);  
  
    function(recv_buf_p);  
}
```

The case of multiple messages from a sender

```
for ( i = 1; i < n_messages ; i++)  
{  
    MPI_Recv(recv_buf_p, recv_buf_size, recv_type,  
            source, MPI_ANY_TAG, comm, status_p);  
  
    function(recv_buf_p);  
}
```

- MPI_ANY_SOURCE and MPI_ANY_TAG are “wildcard” arguments
- Only a receiver can use a wildcard argument.
- Senders must specify a **process rank** and a non-negative **tag**
- No wildcard argument for communicator argument.

The status argument in MPI_Recv

A receiver can receive a message without knowing

- 1) The amount of data in a message
- 2) The sender of the message
- 3) The tag of the message

How can receiver find out these values?

The status argument in MPI_Recv

```
MPI_Status
```

```
{
```

```
  MPI_SOURCE
```

```
  MPI_TAG
```

```
  MPI_ERROR
```

```
  ....
```

```
}
```

```
status.MPI_SOURCE,
```

```
status.MPI_TAG
```

```
MPI_Get_count(&status, recv_type, &count);
```

```
MPI_STATUS_IGNORE
```

Functioning of MPI_Send and MPI_Recv

Sending process will assemble the message “envelope”

1. The sending process will copy the message from send_buf_p to system buffer and attaches the header information containing: sender, receiver, tag, communicator, size
2. Send the message via network switch from the sending process to the receiving process
3. At the receiving end, copy the message from the system buffer to recv_buf_p prescribed by MPI_Recv.

MPI_Send and MPI_Recv are *blocking* and *asynchronous*!

MPI_Send and MPI_Recv behavior

MPI_Send could (1) buffer (2) block

MPI_Recv always blocks

- MPI run time system ensures that the messages be non-overtaking when **two processes** are involved
- However, MPI can't impose performance on a network! when **more than two processes** are involved.

Behavior of Receive Order

```
if (my rank == 0)
{
MPI Send (sendbuf1, count, MPI_INT, 2, tag, comm);
MPI Send (sendbuf2, count, MPI_INT, 1, tag, comm);
}
else if (my rank == 1)
{
MPI Recv (recvbuf1, count, MPI_INT, 0, tag, comm, &status);
MPI Send (recvbuf1, count, MPI_INT, 2, tag, comm);
}
else if (my rank == 2)
{
MPI Recv (recvbuf1, count, MPI_INT, MPI_ANY_SOURCE,
          tag, comm, &status);
MPI Recv (recvbuf2, count, MPI_INT, MPI_ANY_SOURCE,
          tag, comm, &status);
}
```

A few potential pitfalls

Process a	Process b
X	MPI_Recv
MPI_Send	X

Other examples:

1. Non-matching tags
2. Rank of the destination process is the same as the rank of the source process