

MPI: Numerical Integration, P2P and Collective Communication

Kameswararao Anupindi

Department of Mechanical Engineering
Indian Institute of Technology Madras (IITM)

March, 2024

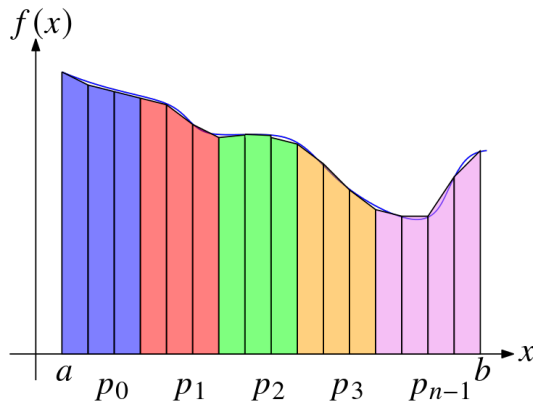


A few potential pitfalls of MPI_Send/MPI_Recv

Process A	Process B
x	MPI_Recv
MPI_Send	x

- ▶ Non-matching tags
- ▶ Rank of the destination process is **the same** as that of the source.

The Trapezoidal Rule approximation



$$\int_a^b f(x)dx = \frac{h}{2} [f(x_0) + f(x_n) + 2 (f(x_1) + f(x_2) \dots + f(x_{n-1}))] \quad (1)$$

The Trapezoidal Rule using MPI in C

```
/* MPI parallel version of trapezoidal rule */  
#include<stdio.h>  
#include<stdlib.h>  
#include<math.h>  
#include<mpi.h>  
  
#define PI 3.14159265358  
  
double func(double x)  
{  
    return (1.0 + sin(x));  
}  
  
double trapezoidal_rule(double la, double lb, double ln, double h)  
{  
    double total;  
    double x;  
    int i;  
  
    total = (func(la) + func(lb))/2.0;  
    for(i = 1; i <= ln-1; i++) /* sharing the work, use only local_n */  
    {  
        x = la + i*h;  
        total += func(x);  
    }  
    total = total * h;  
  
    return total; /* total for each thread, private */  
}
```

The Trapezoidal Rule using MPI in C contd...

```
int main(int argc, char* argv[])
{
    double a, b, final_result, la, lb, lsum, h;
    int myid, nprocs, proc;
    int n, ln;

    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid); /* myrank of the process */
    MPI_Comm_size(MPI_COMM_WORLD, &nprocs); /* size of the communicator */

    n = 1024; /* number of trapezoids.. */
    a = 0.0;
    b = PI; /* hard-coded.. */
    final_result = 0.0;

    h = (b-a)/n;
    ln = n/nprocs; /* nprocs evenly divides number of trapezoids */

    la = a + myid*ln*h;
    lb = la + ln*h;
    lsum = trapezoidal_rule(la, lb, ln, h); /* every process calls this function... */
}
```

The Trapezoidal Rule using MPI in C contd...

```
if (myid != 0)
{
    MPI_Send(&lsum, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD);
}
else /* process 0 */
{
    final_result = lsum;
    for (proc = 1; proc < nprocs; proc++)
    {
        MPI_Recv(&lsum, 1, MPI_DOUBLE, proc, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        final_result += lsum;
    }
}

if (myid == 0) /* output is only printed by process 0 */
{
    printf("\n The area under the curve (1+sin(x)) between 0 to PI is equal to %lf \n\n", final_result);
}

MPI_Finalize();
return 0;
}
```


The Trapezoidal Rule - Enhancements - Dealing with input and output

```
■ if (myid == 0)
{
    printf("\n Enter the lower limit, upper limit and n");
    scanf(&a, &b, &n);

    for (proc = 1, proc<nprocs; proc++)
    {
        MPI_Send(&a, ....);
        MPI_Send(&b, ....);
        MPI_Send(&n, ....);
    }
else
{
    MPI_Recv(&a, 1, ....);
    MPI_Recv(&b, 1, ....);
    MPI_Recv(&n, 1, ....);
}


return 0;
}
```

The Trapezoidal Rule - Enhancements - Calculating global sum



- ▶ Original sum: 7 receives and adds
- ▶ Tree sum: 3 receives and adds
- ▶ if nprocs = 1024, tree sum would do only 10 receives and adds

The Trapezoidal Rule - Calculating global sum - another way



- ▶ Several possibilities exist
- ▶ A method works best for small trees, and another for large trees!
- ▶ A method may work best for system A, and another for system B.
- ▶ MPI provides a **global sum** that works the best in the form of **Collective Communication**.

Collective Communication - MPI_Reduce

```
int MPI_Reduce(  
    void* input_data_p,  
    void* output_data_p,  
    int count,  
    MPI_Datatype datatype,  
    MPI_Op operator,  
    int root,  
    MPI_Comm communicator);
```

```
MPI_Reduce(sendbuf, recvbuf, count, datatype, op, root, comm, ierror)  
TYPE(*), DIMENSION(:), INTENT(IN) :: sendbuf  
TYPE(*), DIMENSION(:) :: recvbuf  
INTEGER, INTENT(IN) :: count, root  
TYPE(MPI_Datatype), INTENT(IN) :: datatype  
TYPE(MPI_Op), INTENT(IN) :: op  
TYPE(MPI_Comm), INTENT(IN) :: comm  
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```


Collective Communication - MPI_Reduce

```
MPI_Reduce(&lsum, &final_result, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
```

```
call MPI_Reduce(lsum, final_result, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD, mpierror);
```

MPI_MAX	MPI_LOR
MPI_MIN	MPI_BAND
MPI_SUM	MPI_BOR
MPI_PROD	MPI_MAXLOC
MPI_LAND	MPI_MINLOC

Collective communication: Reduce



```
MPI_Reduce(  
    void *send_buffer,  
    void *receive_buffer,  
    int count,  
    MPI_Datatype datatype,  
    MPI_Op operator,  
    int root,  
    MPI_Comm communicator)
```

Difference between Collective and P2P communications

- ▶ All the processes must call the same MPI Collective Communication (CC)
- ▶ The arguments passed by each process to MPI CC must be *compatible*
- ▶ All processes must supply an output_data_p, although this is needed only on *root*
- ▶ While P2P are matched using *communicator* and *tags*, MPI CC are matched solely on the basis of *communicator* and order of calling.