

MPI: Numerical Integration, P2P and Collective Communication

Kameswararao Anupindi

Department of Mechanical Engineering
Indian Institute of Technology Madras (IITM)


March, 2024



Other Derived Data Types: Contiguous data

```
int MPI_Type_contiguous(  
    int          count,  
    MPI_Datatype oldtype,  
    MPI_Datatype* newtype);  
  
MPI_TYPE_CONTIGUOUS(count, oldtype, newtype, ierror)
```

Other Derived Data Types: Vector data



- ▶ blocklength =
- ▶ count =
- ▶ stride =

Where would such a pattern of values that has blocks and gaps is needed?

Other Derived Data Types: Vector data

```
int MPI_Type_vector(  
    int          count,  
    int          blocklength,  
    int          stride,  
    MPI_Datatype oldtype,  
    MPI_Datatype* newtype);  
  
MPI_TYPE_VECTOR(count, blocklength, stride, oldtype, newtype, ierror)
```


Representing 2D arrays in C

$$A = \begin{bmatrix} a[0][3] & \dots & \dots & a[3][3] \\ a[0][2] & \dots & \dots & \dots \\ a[0][1] & a[1][1] & \dots & \dots \\ a[0][0] & a[1][0] & a[2][0] & a[3][0] \end{bmatrix}$$

$$A = \begin{bmatrix} 4 & \dots & \dots & 16 \\ 3 & \dots & \dots & \dots \\ 2 & 6 & \dots & \dots \\ 1 & 5 & 9 & 13 \end{bmatrix}$$

1	2	3	4	5										16
---	---	---	---	---	--	--	--	--	--	--	--	--	--	----

Extracting a sub-array in C



- ▶ `blocklength =`
- ▶ `count =`
- ▶ `stride =`

Sending a sub-array

```
MPI_Type_vector(count, blocklength, stride, oldtype, &newtype)
```

```
MPI_Send(&x[1][1], 1, &newtype, ...)
```



Representing 2D arrays in FORTRAN

$$A = \begin{bmatrix} a[1][4] & \dots & \dots & a[4][4] \\ a[1][3] & \dots & \dots & \dots \\ a[1][2] & a[2][2] & \dots & \dots \\ a[1][1] & a[2][1] & a[3][1] & a[4][1] \end{bmatrix}$$

$$A = \begin{bmatrix} 13 & \dots & \dots & 16 \\ 9 & \dots & \dots & \dots \\ 5 & 6 & \dots & \dots \\ 1 & 2 & 3 & 4 \end{bmatrix}$$

1	2	3	4	5										16
---	---	---	---	---	--	--	--	--	--	--	--	--	--	----

Extracting a sub-array in FORTRAN



- ▶ blocklength =
- ▶ count =
- ▶ stride =


Sending a sub-array

```
MPI_Type_vector(count, blocklength, stride, oldtype, &newtype)
```

```
MPI_Send(&x[2][2], 1, &newtype, ...)
```

Remember to commit the new type!

2D or 3D Jacobi/GS



Bubble Sort

```
void Bubble_sort(int a[], int n,)
{
    int list_length, i, temp;

    for (list_length = n; list_length >= 2; list_length--)
        for (i = 0; i < list_length-1; i++)
            if (a[i] > a[i+1])
                {
                    temp = a[i];
                    a[i] = a[i+1];
                    a[i+1] = temp;
                }
}
```

Odd-Even Transposition Sort

```
for (pass = 0; pass < n; pass++)  
    if (pass%2 == 0) {  
        for (i = 1; i < n; i += 2)  
            if (a[i-1] > a[i]) {  
                temp = a[i];  
                a[i] = a[i-1];  
                a[i-1] = temp;  
            }  
    }else {  
        for (i = 1; i < n-1; i += 2)  
            if (a[i] > a[i+1]) {  
                temp = a[i];  
                a[i] = a[i+1];  
                a[i+1] = temp;  
            }  
    }  
}
```

Odd-Even Transposition Sort contd...

Even-pass: $(a[0], a[1]), (a[2], a[3]), (a[4], a[5]), \dots$

Odd-pass: $(a[1], a[2]), (a[3], a[4]), (a[5], a[6]), \dots$

Given list: 5, 9, 4, 3


Even-pass: (5, 9), (4, 3) \rightarrow 5, 9, 3, 4

Odd-pass: 5, (9, 3), 4 \rightarrow 5, 3, 9, 4

Even-pass: (5, 3), (9, 4) \rightarrow 3, 5, 4, 9

Odd-pass: 3, (5, 4), 9 \rightarrow 3, 4, 5, 9

Parallel Odd-Even Transposition sort for $n = p$



Parallel Odd-Even Transposition sort for $n \gg p$

	Process 0	Process 1	Process 2	Process 3
Given	15, 11, 9, 16	3, 14, 8, 7	4, 6, 12, 10	5, 2, 13, 1
After local sort	9, 11, 15, 16	3, 7, 8, 14	4, 6, 10, 12	1, 2, 5, 13
After phase 0	3, 7, 8, 9	11, 14, 15, 16	1, 2, 4, 5	6, 10, 12, 13
After phase 1	3, 7, 8, 9	1, 2, 4, 5	11, 14, 15, 16	6, 10, 12, 13
After phase 2	1, 2, 3, 4	5, 7, 8, 9	6, 10, 11, 12	13, 14, 15, 16
After phase 3	1, 2, 3, 4	5, 6, 7, 8	9, 10, 11, 12	13, 14, 15, 16

Parallel Odd-Even Transposition Sort – Algorithm

```
Sort local elements;
for (pass = 0; pass < comm_sz; pass++) {
    partner = compute_partner(pass, my_rank);
    if (I am active) {
        Send my elements to partner;
        Receive elements from partner;
        if (my_rank < partner)
            Keep smaller elements;
        else
            Keep larger elements;
    }
}
```

Safety in MPI programs

```
MPI_Send(my_elements, n/p, MPI_INT, partner, 0, comm);  
MPI_Recv(temp_elements, n/p, MPI_INT, partner, 0, comm, &status);
```

- ▶ A program that relies on MPI-provided buffering is **unsafe**
- ▶ How can we tell if a program is unsafe?
- ▶ How to modify the communication to make it safe?

How can we tell if a program is safe?

MPI_Send --> MPI_Ssend

```
MPI_Ssend (  
    void*          message_buffer_p,  
    int           message_size,  
    MPI_Datatype   message_type,  
    int           dest_process,  
    int           tag,  
    MPI_Comm       communicator);
```

How to modify the communication to make it safe?

```
MPI_Send(msg, size, MPI_INT, (myid+1)%p, 0, comm);  
MPI_Recv(new_msg, size, MPI_INT, (myid+p-1)%p, 0, comm, &status);
```

```
if (myid % 2 == 0){  
    MPI_Send(msg, size, MPI_INT, (myid+1)%p, 0, comm);  
    MPI_Recv(new_msg, size, MPI_INT, (myid+p-1)%p, 0, comm, &status);  
}  
else{  
    MPI_Recv(new_msg, size, MPI_INT, (myid+p-1)%p, 0, comm, &status);  
    MPI_Send(msg, size, MPI_INT, (myid+1)%p, 0, comm);  
}
```

MPI alternative to manual scheduling

```
int MPI_Sendrecv(  
    void*          send_buf_p,  
    int           send_buf_size,  
    MPI_Datatype   send_buf_type,  
    int           dest,  
    int           send_tag,  
    void*          recv_buf_p,  
    int           recv_buf_size,  
    MPI_Datatype   recv_buf_type,  
    int           source,  
    int           recv_tag,  
    MPI_Comm       communicator,  
    MPI_Status*    status_p);
```

For the same send/recv buffers

```
int MPI_Sendrecv_replace(  
    void*      buf_p,  
    int        buf_size,  
    MPI_Datatype buf_type,  
    int        dest,  
    int        send_tag,  
    int        source,  
    int        recv_tag,  
    MPI_Comm   communicator,  
    MPI_Status* status_p);
```