

# MPI: Numerical Integration, P2P and Collective Communication

Kameswararao Anupindi

Department of Mechanical Engineering  
Indian Institute of Technology Madras (IITM)

March, 2024



1 / 53


## A few potential pitfalls of MPI\_Send/MPI\_Recv

Process A	Process B
x	MPI_Recv
MPI_Send	x

- ▶ Non-matching tags
- ▶ Rank of the destination process is **the same** as that of the source.

2 / 53

## The Trapezoidal Rule approximation



$$\int_a^b f(x)dx = \frac{h}{2} [f(x_0) + f(x_n) + 2(f(x_1) + f(x_2) \dots + f(x_{n-1}))] \quad (1)$$

3/53

## The Trapezoidal Rule using MPI in C

```
/* MPI parallel version of trapezoidal rule */  
#include<stdio.h>  
#include<stdlib.h>  
#include<math.h>  
#include<mpi.h>  
  
#define PI 3.14159265358  
  
double func(double x)  
{  
    return (1.0 + sin(x));  
}  
  
double trapezoidal_rule(double la, double lb, double ln, double h)  
{  
    double total;  
    double x;  
    int i;  
  
    total = (func(la) + func(lb))/2.0;  
    for(i = 1; i <= ln-1; i++) /* sharing the work, use only local_n */  
    {  
        x = la + i*h;  
        total += func(x);  
    }  
    total = total * h;  
  
    return total; /* total for each thread, private */  
}
```

4/53

## The Trapezoidal Rule using MPI in C contd...

```
int main(int argc, char* argv[])
{
    double a, b, final_result, la, lb, lsum, h;
    int myid, nprocs, proc;
    int n, ln;

    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid); /* myrank of the process */
    MPI_Comm_size(MPI_COMM_WORLD, &nprocs); /* size of the communicator */

    n = 1024; /* number of trapezoids.. */
    a = 0.0;
    b = PI; /* hard-coded.. */
    final_result = 0.0;

    h = (b-a)/n;
    ln = n/nprocs; /* nprocs evenly divides number of trapezoids */

    la = a + myid*ln*h;
    lb = la + ln*h;
    lsum = trapezoidal_rule(la, lb, ln, h); /* every process calls this function... */
}
```

5/53

## The Trapezoidal Rule using MPI in C contd...

```
if (myid != 0)
{
    MPI_Send(&lsum, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD);
}
else /* process 0 */
{
    final_result = lsum;
    for (proc = 1; proc < nprocs; proc++)
    {
        MPI_Recv(&lsum, 1, MPI_DOUBLE, proc, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        final_result += lsum;
    }
}

if (myid == 0) /* output is only printed by process 0 */
{
    printf("\n The area under the curve (1+sin(x)) between 0 to PI is equal to %lf \n\n", final_result);
}

MPI_Finalize();
return 0;
}
```

6/53

# The Trapezoidal Rule - Enhancements - Dealing with input and output


```
if (myid == 0)
{
    printf("\n Enter the lower limit, upper limit and n");
    scanf(&a, &b, &n);

    for (proc = 1, proc < nprocs; proc++)
    {
        MPI_Send(&a, ....);
        MPI_Send(&b, ....);
        MPI_Send(&n, ....);
    }
}
else
{
    MPI_Recv(&a, 1, ....);
    MPI_Recv(&b, 1, ....);
    MPI_Recv(&n, 1, ....);
}

return 0;
}
```

7/53

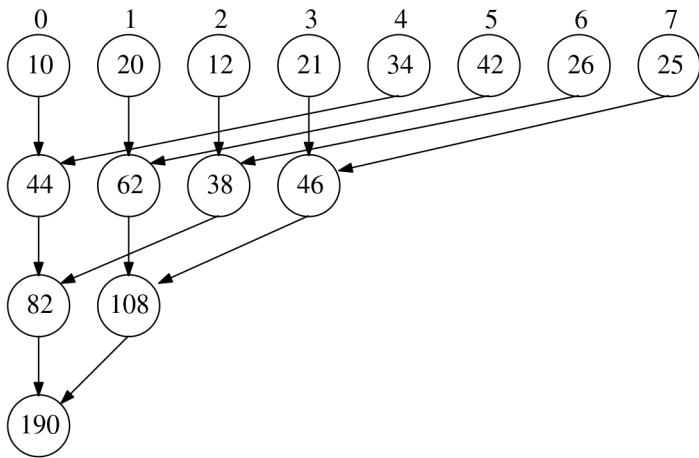
# The Trapezoidal Rule - Enhancements - Calculating global sum



- ▶ Original sum: 7 receives and adds
- ▶ Tree sum: 3 receives and adds
- ▶ if nprocs = 1024, tree sum would do only 10 receives and adds

8/53

## The Trapezoidal Rule - Calculating global sum - another way



- ▶ Several possibilities exist
- ▶ A method works best for small trees, and another for large trees!
- ▶ A method may work best for system A, and another for system B.
- ▶ MPI provides a **global sum** that works the best in the form of **Collective Communication**.

9/53

## Collective Communication - MPI\_Reduce

```
int MPI_Reduce(  
    void* input_data_p,  
    void* output_data_p,  
    int count,  
    MPI_Datatype datatype,  
    MPI_Op operator,  
    int root,  
    MPI_Comm communicator);
```

```
MPI_Reduce(sendbuf, recvbuf, count, datatype, op, root, comm, ierror)  
TYPE(*), DIMENSION(:), INTENT(IN) :: sendbuf  
TYPE(*), DIMENSION(:) :: recvbuf  
INTEGER, INTENT(IN) :: count, root  
TYPE(MPI_Datatype), INTENT(IN) :: datatype  
TYPE(MPI_Op), INTENT(IN) :: op  
TYPE(MPI_Comm), INTENT(IN) :: comm  
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

10/53


# Collective Communication - MPI\_Reduce

```
MPI_Reduce(&lsum, &final_result, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
```

```
call MPI_Reduce(lsum, final_result, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD, mpierror);
```

MPI_MAX	MPI_LOR
MPI_MIN	MPI_BAND
MPI_SUM	MPI_BOR
MPI_PROD	MPI_MAXLOC
MPI_LAND	MPI_MINLOC

## Collective communication: Reduce



```
MPI_Reduce(  
    void *send_buffer,  
    void *receive_buffer,  
    int count,  
    MPI_Datatype datatype,  
    MPI_Op operator,  
    int root,  
    MPI_Comm communicator)
```

## Difference between Collective and P2P communications

- ▶ All the processes must call the same MPI Collective Communication (CC)
- ▶ The arguments passed by each process to MPI CC must be *compatible*
- ▶ All processes must supply an *output\_data\_p*, although this is needed only on *root*
- ▶ While P2P are matched using *communicator* and *tags*, MPI CC are matched solely on the basis of *communicator* and order of calling.

13/53

## Multiple CC calls

```
Process 0:
  a = 1, b = 0, c = 2, d = 0;
  dest_process = 0;

  MPI_Reduce(&a, &b, ..., 0, comm);
  MPI_Reduce(&c, &d, ..., 0, comm);

Process 1:
  a = 1, b = 0, c = 2, d = 0;
  dest_process = 0;

  MPI_Reduce(&c, &d, ..., 0, comm);
  MPI_Reduce(&a, &b, ..., 0, comm);

Process 2:
  a = 1, b = 0, c = 2, d = 0;
  dest_process = 0;

  MPI_Reduce(&a, &b, ..., 0, comm);
  MPI_Reduce(&c, &d, ..., 0, comm);
```

14/53

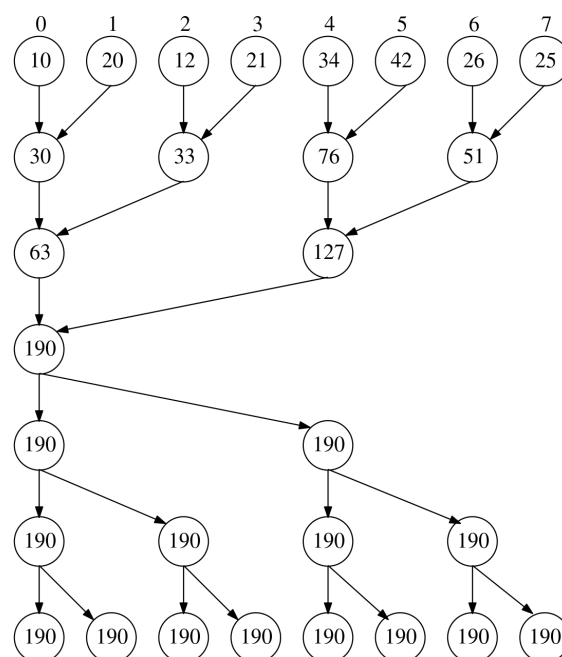
## Reduction on the same variable

```
MPI_Reduce(&x, &x, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);
```

- ▶ Illegal in MPI
- ▶ Produces unpredictable result.

15/53


## MPI\_Allreduce: Tree and Reverse-tree



16/53



## MPI\_Allreduce: Butterfly



17/53


## MPI\_Allreduce function prototype

```
int MPI_Allreduce(  
    void* input_data_p,  
    void* output_data_p,  
    int count,  
    MPI_Datatype datatype,  
    MPI_Op operator,  
    MPI_Comm communicator);
```

```
MPI_ALLREDUCE(sendbuf, recvbuf, count, datatype, op, comm, ierror)  
TYPE(*), DIMENSION(:), INTENT(IN) :: sendbuf  
TYPE(*), DIMENSION(:) :: recvbuf  
INTEGER, INTENT(IN) :: count  
TYPE(MPI_Datatype), INTENT(IN) :: datatype  
TYPE(MPI_Op), INTENT(IN) :: op  
TYPE(MPI_Comm), INTENT(IN) :: comm  
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

18/53

## Collective communication: Broadcast




```
MPI_Bcast(  
    void *data,  
    int count,  
    MPI_Datatype datatype,  
    int root,  
    MPI_Comm communicator)
```

- ▶ Use a tree-structured communication instead!
- ▶ data\_p is an input argument on root (send\_proc) and output on the other processes.

19/53


## Collective communication: Scatter



```
MPI_Scatter(  
    void *send_data,  
    int send_count,  
    MPI_Datatype datatype,  
    void *receive_data,  
    int receive_count,  
    MPI_Datatype datatype,  
    int root,  
    MPI_Comm communicator)
```

20/53

## Collective communication: Gather



```
MPI_Gather(  
    void *send_data,  
    int send_count,  
    MPI_Datatype datatype,  
    void *receive_data,  
    int receive_count,  
    MPI_Datatype datatype,  
    int root,  
    MPI_Comm communicator)
```

21/53

## Broadcast example program

```
if (myid == 0)  
    buf = 327;  
  
MPI_Bcast(&buf, 1, MPI_INT, 0, MPI_COMM_WORLD);  
  
if (myid == 0)  
    printf("\n Broadcasted values on processors are:\n");  
  
printf("\t (%d, %d)\n", myid, buf);
```

22/53

## Gather example program

```
int send_buf, *recv_buf;
if (myid == 0)
{
    recv_buf = (int *)malloc(size*sizeof(int));
}

send_buf = 100+myid*myid;

MPI_Gather(&send_buf, 1, MPI_INT, recv_buf, 1, MPI_INT, 0, MPI_COMM_WORLD);

if (myid == 0)
{
    printf("\n Received values on host process are:\n");
    for(i=0; i<size; i++)
        printf("\t %d", recv_buf[i]);
    printf("\n");
}

if (myid == 0)
    free(recv_buf);
```

23/53

## Scatter example program

```
int *send_buf, recv_buf;
if (myid == 0)
{
    send_buf = (int *)malloc(size*sizeof(int));
    for(i=0; i<size; i++)
        send_buf[i] = 100+i*5+i;
}

MPI_Scatter(send_buf, 1, MPI_INT, &recv_buf, 1, MPI_INT, 0, MPI_COMM_WORLD);

if (myid == 0)
    printf("\n Received values on processors are:\n");

printf("\t (%d, %d)", myid, recv_buf);

if (myid == 0)
    free(send_buf);
```

24/53

## Matrix - Vector Multiplication using block decomposition

$$\begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \cdots & a_{nn} \end{bmatrix} \times \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{bmatrix}$$

- ▶ Row block-decomposition
- ▶ Consider

$$\begin{bmatrix} x_1 & x_2 & \cdots & x_n \end{bmatrix}^T$$

has block-decomposition as well

- ▶ How to arrange that each process has access to all components of  $[x]$ ?

25/53

## MPI\_Allgather

```
int MPI_Allgather(  
    const void* send_buf_p,  
    int send_count,  
    MPI_Datatype send_type,  
    void* recv_buf_p,  
    int recv_count,  
    MPI_Datatype recv_type,  
    MPI_Comm communicator);
```

```
MPI_Allgather(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvttype, comm, ierror)  
TYPE(*), DIMENSION(..), INTENT(IN) :: sendbuf  
TYPE(*), DIMENSION(..) :: recvbuf  
INTEGER, INTENT(IN) :: sendcount, recvcount  
TYPE(MPI_Datatype), INTENT(IN) :: sendtype, recvttype  
TYPE(MPI_Comm), INTENT(IN) :: comm  
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

26/53

## Motivation for MPI Derived Datatypes

```
double x[1000];

if (myid == 0)
    for(i = 0; i < 1000; i++)
        MPI_Send(&x[i], 1, MPI_DOUBLE, 1, 0, comm);
else
    for(i = 0; i < 1000; i++)
        MPI_Recv(&x[i], 1, MPI_DOUBLE, 0, 0, comm, &status);
```

```
double x[1000];

if (myid == 0)
    MPI_Send(x, 1000, MPI_DOUBLE, 1, 0, comm);
else
    MPI_Recv(x, 1000, MPI_DOUBLE, 0, 0, comm, &status);
```

27/53

## MPI Data consolidation

- ▶ The **count** argument
- ▶ MPI Derived Datatypes
- ▶ MPI\_Pack and MPI\_Unpack

28/53

# MPI Derived Datatypes

- Collection of data items in memory by storing both **type** and **relative memory locations**

Variable	Address
a	24
b	40
c	48

```
{(MPI_DOUBLE,0), (MPI_DOUBLE,16), (MPI_DOUBLE,24)}
```

29/53

## Building MPI Derived Datatypes

```
int MPI_Type_create_struct(  
    int          count,  
    int          array_of_blocklengths[],  
    MPI_Aint     array_of_displacements[],  
    MPI_Datatype array_of_types[],  
    MPI_Datatype* new_type_p);  
  
int array_of_blocklengths[3] = {1, 1, 1};  
  
array_of_blocklengths[0] = 5;  
  
array_of_displacements[] = {0, 16, 24};
```

30/53

## Building MPI Derived Datatypes contd...

```
int MPI_Get_address(
    void*      location_p,
    MPI_Aint*   address_p);

MPI_Aint a_addr, b_addr, c_addr;

MPI_Get_address(&a, &a_addr);
array_of_displacements[0] = 0;

MPI_Get_address(&b, &b_addr);
array_of_displacements[1] = b_addr - a_addr;

MPI_Get_address(&c, &c_addr);
array_of_displacements[2] = c_addr - a_addr;
```

31/53

## Building MPI Derived Datatypes contd...

```
MPI_Datatype array_of_types[3] = {MPI_DOUBLE, MPI_DOUBLE, MPI_INT};
....
....
MPI_Datatype new_mpi_t;
...
MPI_Type_create_struct(3, array_of_blocklengths,
    array_of_displacements, array_of_types, &new_mpi_t);
...
int MPI_Type_commit(MPI_Datatype* new_mpi_type_p);
...
...
MPI_Bcast(&a, 1, new_mpi_t, 0, comm);
...
...
int MPI_Type_free(MPI_Datatype* old_mpi_type_p);
```

32/53




## Other Derived Data Types: Contiguous data

```
int MPI_Type_contiguous(  
    int          count,  
    MPI_Datatype oldtype,  
    MPI_Datatype* newtype);  
  
MPI_TYPE_CONTIGUOUS(count, oldtype, newtype, ierror)
```

33 / 53

## Other Derived Data Types: Vector data



- ▶ blocklength =
- ▶ count =
- ▶ stride =

Where would such a pattern of values that has blocks and gaps is needed?

34 / 53

# Other Derived Data Types: Vector data


```
int MPI_Type_vector(  
    int          count,  
    int          blocklength,  
    int          stride,  
    MPI_Datatype oldtype,  
    MPI_Datatype* newtype);  
  
MPI_TYPE_VECTOR(count, blocklength, stride, oldtype, newtype, ierror)
```

# Representing 2D arrays in C

$$A = \begin{bmatrix} a[0][3] & \dots & \dots & a[3][3] \\ a[0][2] & \dots & \dots & \dots \\ a[0][1] & a[1][1] & \dots & \dots \\ a[0][0] & a[1][0] & a[2][0] & a[3][0] \end{bmatrix}$$
$$A = \begin{bmatrix} 4 & \dots & \dots & 16 \\ 3 & \dots & \dots & \dots \\ 2 & 6 & \dots & \dots \\ 1 & 5 & 9 & 13 \end{bmatrix}$$

1	2	3	4	5											16
---	---	---	---	---	--	--	--	--	--	--	--	--	--	--	----

## Extracting a sub-array in C



- ▶ blocklength =
- ▶ count =
- ▶ stride =

37/53

## Sending a sub-array

```
MPI_Type_vector(count, blocklength, stride, oldtype, &newtype)  
MPI_Send(&x[1][1], 1, &newtype, ...)
```

38/53

# Representing 2D arrays in FORTRAN

$$A = \begin{bmatrix} a[1][4] & \dots & \dots & a[4][4] \\ a[1][3] & \dots & \dots & \dots \\ a[1][2] & a[2][2] & \dots & \dots \\ a[1][1] & a[2][1] & a[3][1] & a[4][1] \end{bmatrix}$$

$$A = \begin{bmatrix} 13 & \dots & \dots & 16 \\ 9 & \dots & \dots & \dots \\ 5 & 6 & \dots & \dots \\ 1 & 2 & 3 & 4 \end{bmatrix}$$

1	2	3	4	5											16
---	---	---	---	---	--	--	--	--	--	--	--	--	--	--	----

39/53

# Extracting a sub-array in FORTRAN


--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

- ▶ blocklength =
- ▶ count =
- ▶ stride =

40/53

## Sending a sub-array


```
MPI_Type_vector(count, blocklength, stride, oldtype, &newtype)
```

```
MPI_Send(&x[2][2], 1, &newtype, ...)
```

Remember to commit the new type!

41/53

## 2D or 3D Jacobi/GS



42/53

## Bubble Sort

```
void Bubble_sort(int a[], int n,)
{
    int list_length, i, temp;

    for (list_length = n; list_length >= 2; list_length--)
        for (i = 0; i < list_length-1; i++)
            if (a[i] > a[i+1])
                {
                    temp = a[i];
                    a[i] = a[i+1];
                    a[i+1] = temp;
                }
}
```

43/53

## Odd-Even Transposition Sort

```
for (pass = 0; pass < n; pass++)
    if (pass%2 == 0) {
        for (i = 1; i < n; i += 2)
            if (a[i-1] > a[i]) {
                temp = a[i];
                a[i] = a[i-1];
                a[i-1] = temp;
            }
    }else {
        for (i = 1; i < n-1; i += 2)
            if (a[i] > a[i+1]) {
                temp = a[i];
                a[i] = a[i+1];
                a[i+1] = temp;
            }
    }
}
```

44/53

# Odd-Even Transposition Sort contd...


Even-pass: (a[0], a[1]), (a[2], a[3]), (a[4], a[5]), ...

Odd-pass: (a[1], a[2]), (a[3], a[4]), (a[5], a[6]), ...

Given list: 5, 9, 4, 3  
Even-pass: (5, 9), (4, 3) -> 5, 9, 3, 4  
Odd-pass: 5, (9, 3), 4 -> 5, 3, 9, 4  
Even-pass: (5, 3), (9, 4) -> 3, 5, 4, 9  
Odd-pass: 3, (5, 4), 9 -> 3, 4, 5, 9

45/53

## Parallel Odd-Even Transposition sort for $n = p$



46/53

## Parallel Odd-Even Transposition sort for $n \gg p$

	Process 0	Process 1	Process 2	Process 3
Given	15, 11, 9, 16	3, 14, 8, 7	4, 6, 12, 10	5, 2, 13, 1
After local sort	9, 11, 15, 16	3, 7, 8, 14	4, 6, 10, 12	1, 2, 5, 13
After phase 0	3, 7, 8, 9	11, 14, 15, 16	1, 2, 4, 5	6, 10, 12, 13
After phase 1	3, 7, 8, 9	1, 2, 4, 5	11, 14, 15, 16	6, 10, 12, 13
After phase 2	1, 2, 3, 4	5, 7, 8, 9	6, 10, 11, 12	13, 14, 15, 16
After phase 3	1, 2, 3, 4	5, 6, 7, 8	9, 10, 11, 12	13, 14, 15, 16

47/53

## Parallel Odd-Even Transposition Sort – Algorithm

```
Sort local elements;
for (pass = 0; pass < comm_sz; pass++) {
    partner = compute_partner(pass, my_rank);
    if (I am active) {
        Send my elements to partner;
        Receive elements from partner;
        if (my_rank < partner)
            Keep smaller elements;
        else
            Keep larger elements;
    }
}
```

48/53



## Safety in MPI programs

```
MPI_Send(my_elements, n/p, MPI_INT, partner, 0, comm);  
MPI_Recv(temp_elements, n/p, MPI_INT, partner, 0, comm, &status);
```

- ▶ A program that relies on MPI-provided buffering is **unsafe**
- ▶ How can we tell if a program is unsafe?
- ▶ How to modify the communication to make it safe?

49/53

## How can we tell if a program is safe?

```
MPI_Send --> MPI_Ssend
```

```
MPI_Ssend (  
    void*          message_buffer_p,  
    int            message_size,  
    MPI_Datatype    message_type,  
    int            dest_process,  
    int            tag,  
    MPI_Comm        communicator);
```

50/53

## How to modify the communication to make it safe?

```
MPI_Send(msg, size, MPI_INT, (myid+1)%p, 0, comm);  
MPI_Recv(new_msg, size, MPI_INT, (myid+p-1)%p, 0, comm, &status);
```

```
if (myid % 2 == 0){  
    MPI_Send(msg, size, MPI_INT, (myid+1)%p, 0, comm);  
    MPI_Recv(new_msg, size, MPI_INT, (myid+p-1)%p, 0, comm, &status);  
}  
else{  
    MPI_Recv(new_msg, size, MPI_INT, (myid+p-1)%p, 0, comm, &status);  
    MPI_Send(msg, size, MPI_INT, (myid+1)%p, 0, comm);  
}
```

51/53

## MPI alternative to manual scheduling

```
int MPI_Sendrecv(  
    void*          send_buf_p,  
    int            send_buf_size,  
    MPI_Datatype   send_buf_type,  
    int            dest,  
    int            send_tag,  
    void*          recv_buf_p,  
    int            recv_buf_size,  
    MPI_Datatype   recv_buf_type,  
    int            source,  
    int            recv_tag,  
    MPI_Comm       communicator,  
    MPI_Status*    status_p);
```

52/53

## For the same send/recv buffers

```
int MPI_Sendrecv_replace(
    void*      buf_p,
    int        buf_size,
    MPI_Datatype buf_type,
    int        dest,
    int        send_tag,
    int        source,
    int        recv_tag,
    MPI_Comm   communicator,
    MPI_Status* status_p);
```